

Theoretical Key Recovery from Truncated Differentials: A BCT-Algebraic Wrong-Key Profile for SPECK32/64

Anonymous Author(s)

Abstract. Neural-network-based distinguishers have achieved state-of-the-art results in differential cryptanalysis of SPECK32/64, most notably through Gohr’s CRYPTO 2019 attack framework. That framework rests on two black-box components—a deep residual distinguisher and an empirically sampled wrong-key response profile—whose internal structure has so far resisted closed-form analysis. In this paper we replace both components with algebraically defined counterparts and thereby provide a fully interpretable reconstruction of the attack. We substitute the neural distinguisher with a truncated differential distinguisher whose scores are computed directly from the cipher’s difference distribution table (DDT), and we derive the wrong-key profile analytically via a novel application of the boomerang connectivity table (BCT). Specifically, for each 16-bit subkey offset g we show that the expected distinguisher score under the corresponding wrong-key hypothesis equals a *truncated self-returning BCT probability* $p_{\text{trunc}}^{\text{self}}(g)$, which admits an exact bit-by-bit computation through a carry-propagation transfer matrix. The resulting profile is derived rather than measured, its shape is governed by the algebraic structure of SPECK32/64’s modular-addition round function, and it requires no neural network training or large-scale sampling. We demonstrate that the reconstructed Bayesian beam-search framework achieves key-recovery performance competitive with Gohr’s original attack while being amenable to complexity analysis. Our results suggest that the statistical learning structure underlying neural-distinguisher attacks on ARX ciphers can in many cases be recovered analytically from classical differential tools.

Keywords: Speck32/64 · differential cryptanalysis · truncated differentials · boomerang connectivity table · wrong-key profile · key recovery · ARX cipher

1 Introduction

Differential cryptanalysis, introduced by Biham and Shamir [6], exploits the non-uniform propagation of XOR differences through a block cipher’s round function to recover secret key material. For an n -bit block cipher, the fundamental object governing this propagation is the *difference distribution table* (DDT): an array indexed by input and output differences whose entries record the probability

that a uniformly random input pair with the given input difference produces the given output difference after one round of encryption. The boomerang attack, introduced by Wagner [17] and subsequently refined through the notion of the *boomerang connectivity table* (BCT) by Cid et al. [7], extends the differential framework to cover compatible transitions across the encryption and decryption directions. Both tables have become standard tools in the security evaluation of symmetric-key primitives.

Speck32/64 and classical attacks. The SPECK32/64 cipher [4] belongs to the ARX (Add-Rotate-XOR) design family and has attracted sustained cryptanalytic interest as a lightweight benchmark target. Its round function consists of a modular addition, a rotation, and an XOR, combined with a key injection:

$$(x, y) \mapsto ((x \ggg 7) \boxplus y \oplus k, (y \lll 2) \oplus ((x \ggg 7) \boxplus y \oplus k)), \quad (1)$$

where $\boxplus$ denotes addition modulo 2^{16} and k is the current round subkey. The differential behaviour of modular addition is well understood [13], enabling precise computation of DDT and BCT entries for individual rounds of SPECK32/64. Classical differential attacks have been mounted on up to 8 rounds of SPECK32/64 using optimised distinguishers [1,11], but extending these attacks further has proven difficult due to the rapid diffusion of differences across rounds.

Gohr’s neural-network attack. In a landmark result at CRYPTO 2019, Gohr [9] demonstrated that deep residual networks could serve as highly effective distinguishers for 7- and 8-round SPECK32/64, and embedded them within a sophisticated Bayesian key-recovery framework to mount practical attacks on up to 11 rounds. The framework integrates three tightly coupled components.

1. **Neural distinguisher.** A deep residual network is trained to classify ciphertext pairs produced under a fixed input difference from randomly generated pairs. Given a ciphertext pair (C, C') with $C \oplus C' = \Delta_{\text{out}}$, the network outputs a probability $p \in (0, 1)$ that the pair arises from a real differential rather than from two independent random plaintexts. The per-pair log-likelihood ratio (LLR) $\log_2(p/(1-p))$ is accumulated over all available ciphertext pairs to produce a score for each subkey hypothesis.
2. **Wrong-key response profile (WKP).** For each of the 2^{16} possible last-round subkey offsets $g = k_r \oplus k_g$ (where k_r is the real subkey and k_g is the guessed subkey), Gohr measures the *mean* and *standard deviation* of the network’s output when ciphertext pairs are partially decrypted under the wrong subkey k_g . This profiling step requires on the order of 10^6 forward passes through the network for each of the 2^{16} offset values, making it computationally intensive and inherently tied to the specific trained network.
3. **Bayesian beam search.** Given a set of ciphertext pairs, Gohr iterates a beam search over the subkey space. In each iteration, the current pool of $B = 32$ subkey candidates is partially decrypted and fed to the neural network; the

resulting empirical output means are compared to the precomputed WKP via a Bayesian distance metric (an ℓ^2 residual against the profile), and the B candidates whose profile best matches the observed statistics are retained for the next iteration. A two-layer nesting recovers the last two round subkeys sequentially, followed by a local refinement step that exhaustively explores candidates within Hamming distance 2 of the current best guess.

This framework achieved results far beyond what had been attained with classical methods and demonstrated, for the first time, that machine-learning-based distinguishers could be integrated into a full key-recovery attack with practical data and time complexity.

The interpretability gap. Despite its empirical success, Gohr’s framework exhibits a fundamental limitation: it is constructed almost entirely from observations rather than proofs. The neural distinguisher is a black box whose internal representations have no direct cryptanalytic interpretation. The WKP is produced by brute-force sampling and carries no closed-form description. As a consequence:

- No theoretical bound governs the *shape* of the WKP, so the signal-to-noise ratio of the beam search cannot be predicted analytically.
- No complexity formula characterises the minimum number of ciphertext pairs required for the search to succeed with a given probability.
- Generalising the attack to a related cipher, or to a different number of rounds, requires repeating the same expensive empirical construction—training a new network and resampling the profile—from scratch.

These limitations have prompted a line of follow-up work aimed at understanding why neural distinguishers work as well as they do [5,3,8,14], but a complete analytical reconstruction of Gohr’s full key-recovery framework has not yet been achieved.

This work: a BCT-algebraic reconstruction. We propose a theoretically grounded reconstruction of Gohr’s key recovery framework in which both black-box components are replaced by algebraically defined counterparts derived from the differential structure of SPECK32/64.

Truncated differential distinguisher. We replace the neural network with a *truncated differential distinguisher*. Rather than learning patterns from training data, the distinguisher assigns to each ciphertext pair a LLR computed directly from the DDT of reduced-round SPECK32/64: after partially decrypting the pair by one round under the candidate subkey, we evaluate the log-probability of the resulting truncated output difference under the target differential distribution. The truncation—fixing only a subset of the 32 output-difference bits—allows us to balance statistical power against computation: a coarser truncation reduces the

distinguisher’s information content but yields a simpler, smaller-footprint probability table. Every score is an explicit function of the observed output difference and the precomputed differential probabilities, with no black-box inference involved.

BCT-algebraic wrong-key profile. The central analytical contribution of this work is a closed-form derivation of the WKP. We show that, for a one-round truncated differential distinguisher applied to SPECK32/64, the expected LLR score under wrong-key hypothesis $g = k_r \oplus k_g$ is determined by the probability that the truncated differential *self-returns* under key offset g :

$$p_{\text{trunc}}^{\text{self}}(g) := \Pr_x[\Delta_{\text{out}} \in \mathcal{T} \mid \Delta_{\text{in}} \in \mathcal{T}, \text{ one round of SPECK32/64 with key offset } g], \quad (2)$$

where $\mathcal{T}$ is the truncated class defined by the chosen mask. We show that $p_{\text{trunc}}^{\text{self}}(g)$ can be computed *exactly* for every $g \in \{0, 1\}^{16}$ via a BCT-style algebraic transfer-matrix computation that processes each bit position independently and propagates carry distributions through a small 16×16 stochastic matrix M . The resulting profile requires no network training and no sampling; it is a direct consequence of the algebraic structure of the ARX round function. In particular, the non-trivial variation of $p_{\text{trunc}}^{\text{self}}(g)$ across key offsets—the phenomenon that makes the Bayesian beam search effective—is explained entirely by the differential properties of modular addition and the specific bit-pattern of the chosen truncated differential.

Interpretable beam-search analysis. With both components in closed form, we provide a theoretical characterisation of the Bayesian beam search. The signal available to the search is the gap $p_{\text{trunc}}^{\text{self}}(0) - p_{\text{trunc}}^{\text{self}}(g)$ between the correct-key and wrong-key self-returning probabilities; its magnitude and distribution over g govern the convergence rate. We derive a lower bound on the expected score advantage of the correct subkey, and use it to bound the data complexity of the attack. We further show that the UCB-guided structure selection adopted by Gohr can be understood as a multi-armed bandit operating on noisy estimates of these same self-returning probabilities.

Contributions. We summarise the contributions of this paper as follows.

1. We introduce the *truncated self-returning BCT probability* $p_{\text{trunc}}^{\text{self}}(g)$ and show that it constitutes the wrong-key profile of any one-round truncated differential distinguisher applied to an ARX cipher whose round function contains a single modular addition (Sect. 3).
2. We present an exact bit-serial transfer-matrix algorithm for computing $p_{\text{trunc}}^{\text{self}}(g)$ for all $g \in \{0, 1\}^{16}$ in time $O(16 \cdot 2^{16})$, requiring no sampling (Sect. 4).
3. We instantiate the framework on SPECK32/64 with two carefully selected truncated differentials (with input differences `0x181E` and `0xC6A5`) and verify that the computed profiles match empirical measurements to within statistical noise (Sect. 7).

4. We integrate the truncated distinguisher and the algebraic WKP into a full reconstruction of Gohr’s Bayesian key-recovery attack and demonstrate key recovery competitive with the original neural-network-based framework on 9–11 rounds of SPECK32/64 (Sect. 5).
5. We derive data-complexity bounds for the beam search in terms of $p_{\text{trunc}}^{\text{self}}$ and compare them to Gohr’s empirically observed data usage (Sect. 6).

Related work. Neural distinguishers for block ciphers have been studied extensively since Gohr’s breakthrough. Benamira et al. [5] gave a partial theoretical explanation of why neural networks can distinguish r -round SPECK32/64 by identifying specific statistical features learned by the network, and showed that the neural distinguisher effectively approximates the cipher’s DDT together with penultimate-round differential distributions. Hou et al. [8] improved distinguisher accuracy through enhanced training procedures. Liu et al. [14] revisited the connection between neural distinguishers and the DDT.

Most directly related to our work is Bao, Lu, Yao, and Zhang [3], who provide an explicit cryptanalytic account of the additional information neural distinguishers exploit beyond the DDT: the correlation between ciphertext values and intermediate-state differences at the last modular addition, recoverable from conditional differential probabilities. Their work significantly advances the interpretability of Gohr’s framework and extends key-recovery to 14 rounds of SPECK32/64 in the related-key setting. Our contribution is orthogonal: rather than explaining what the neural distinguisher *learns*, we replace it entirely with a truncated differential distinguisher and derive its wrong-key profile algebraically via the BCT, without any neural network training or sampling. Together, the two lines of work suggest that the statistical learning structure of Gohr’s framework is recoverable from classical differential tools.

The wrong-key profile as a function of the key offset was studied systematically by Leander and Bao [12], who characterised the non-conformity of the wrong-key randomisation hypothesis and showed how it can be exploited to reduce attack complexity. Our BCT-algebraic formula for $p_{\text{trunc}}^{\text{self}}(g)$ gives a closed-form instance of the same phenomenon for the specific case of a one-round truncated differential on an ARX cipher. On the BCT side, the original formulation [7] and its extensions [16,18] focus on the probability of boomerang quartets in the middle rounds of a cipher; our use of BCT-style algebra to compute single-round wrong-key profiles is, to the best of our knowledge, new. Truncated differentials were introduced by Knudsen [10] and have been applied to SPECK32/64 by several works [1,11]; here we combine them with the Bayesian key-ranking apparatus of Gohr’s framework for the first time.

Organisation. Section 2 reviews SPECK32/64, the DDT, the BCT, and Gohr’s original framework in the notation used throughout the paper. Section 3 introduces the truncated self-returning BCT probability and proves its role as the wrong-key profile. Section 4 presents the transfer-matrix algorithm for computing $p_{\text{trunc}}^{\text{self}}$. Section 5 describes the full reconstructed key-recovery attack. Sec-

tion 6 derives data-complexity bounds. Section 7 presents experimental results. Section 8 concludes and discusses open problems.

2 Preliminaries

2.1 Notation

We write $\mathbb{F}_2^n$ for the vector space of n -bit strings over $\{0, 1\}$ and $\mathbb{Z}_{2^n}$ for the ring of integers modulo 2^n . Throughout the paper the block operates on pairs of 16-bit words; we identify $\mathbb{F}_2^{16}$ with $\mathbb{Z}_{2^{16}}$ whenever the interpretation is clear from context. The XOR of two bit strings $a, b \in \mathbb{F}_2^n$ is written $a \oplus b$, modular addition in $\mathbb{Z}_{2^{16}}$ is written $a \boxplus b$, and bitwise AND is written $a \& b$. For $a \in \mathbb{F}_2^{16}$ we write $a \ggg r$ (resp. $a \lll r$) for the right (resp. left) rotation of a by r bit positions. The Hamming weight of a bit string a is written $\text{wt}(a)$.

2.2 The SPECK32/64 Block Cipher

SPECK32/64 [4] is a member of the Simon/Speck family of lightweight ARX (Add-Rotate-XOR) block ciphers. It operates on a 32-bit block, represented as a pair $(x, y) \in \mathbb{F}_2^{16} \times \mathbb{F}_2^{16}$, under a 64-bit key.

Round function. The encryption round function $F_k : \mathbb{F}_2^{16} \times \mathbb{F}_2^{16} \rightarrow \mathbb{F}_2^{16} \times \mathbb{F}_2^{16}$ parameterised by the round subkey $k \in \mathbb{F}_2^{16}$ is

$$F_k(x, y) = (x', y'), \quad x' = (x \ggg 7) \boxplus y \oplus k, \quad y' = (y \lll 2) \oplus x'. \quad (3)$$

The decryption round function $F_k^{-1}(x', y')$ is given by

$$y = ((y' \oplus x') \ggg 2), \quad x = ((x' \oplus k \ominus y) \lll 7), \quad (4)$$

where $\ominus$ denotes subtraction modulo 2^{16} .

Key schedule. A 64-bit master key is parsed as four 16-bit words $(k^{(3)}, k^{(2)}, k^{(1)}, k^{(0)})$. The initial subkey is $k_0 = k^{(3)}$, and three auxiliary words $(\ell_0, \ell_1, \ell_2) = (k^{(0)}, k^{(1)}, k^{(2)})$ are maintained. For $i = 0, 1, \dots, r - 2$ (with r the number of rounds):

$$\ell_{i+3} = (\ell_i \ggg 7) \boxplus k_i \oplus i, \quad k_{i+1} = (k_i \lll 2) \oplus \ell_{i+3}. \quad (5)$$

The full key schedule for $r = 22$ rounds (the full SPECK32/64 specification) produces subkeys $k_0, k_1, \dots, k_{21}$. In our attack we target a reduced-round variant with $r \in \{7, 8, \dots, 11\}$ rounds; the last round subkey is k_{r-1} and the penultimate is k_{r-2} .

2.3 Differential Cryptanalysis and the DDT

Let $E_k : \mathbb{F}_2^{32} \rightarrow \mathbb{F}_2^{32}$ denote an r -round block cipher with key k . Given a pair of plaintexts (P, P') with XOR difference $\Delta_{\text{in}} = P \oplus P'$, the corresponding output difference is $\Delta_{\text{out}} = E_k(P) \oplus E_k(P')$. A *differential* $(\Delta_{\text{in}} \rightarrow \Delta_{\text{out}})$ holds with *differential probability*

$$\text{DP}(\Delta_{\text{in}} \rightarrow \Delta_{\text{out}}) = \Pr_P[E_k(P) \oplus E_k(P \oplus \Delta_{\text{in}}) = \Delta_{\text{out}}]. \quad (6)$$

The *difference distribution table* (DDT) of a mapping $f : \mathbb{F}_2^n \rightarrow \mathbb{F}_2^m$ is the $2^n \times 2^m$ array

$$\text{DDT}[a][b] = \#\{x \in \mathbb{F}_2^n \mid f(x) \oplus f(x \oplus a) = b\}, \quad (7)$$

so that $\text{DP}(a \rightarrow b) = \text{DDT}[a][b]/2^n$ for a uniform input distribution. For the modular-addition map $\text{add}(x, y) = x \boxplus y$, an efficient bit-parallel formula for $\text{DDT}[(\delta_x, \delta_y)][\delta_z]$ was given by Lipmaa and Moriai [13] and provides the computational backbone for DDT-based distinguishers on ARX ciphers.

2.4 Truncated Differentials

Truncated differentials, introduced by Knudsen [10], relax the requirement that a specific output difference is observed, instead specifying only that the output difference falls within some set $\mathcal{T} \subseteq \mathbb{F}_2^{32}$.

Definition 1 (Truncated differential class). Let $\alpha, \mu_\alpha \in \mathbb{F}_2^{16}$ and $\beta, \mu_\beta \in \mathbb{F}_2^{16}$ be representative and free-bit mask words, respectively. The associated truncated class is

$$\mathcal{T}(\alpha, \mu_\alpha, \beta, \mu_\beta) := \{(\Delta_L, \Delta_R) \in \mathbb{F}_2^{16} \times \mathbb{F}_2^{16} \mid \Delta_L \& \overline{\mu_\alpha} = \alpha \& \overline{\mu_\alpha}, \Delta_R \& \overline{\mu_\beta} = \beta \& \overline{\mu_\beta}\}, \quad (8)$$

where $\overline{\mu}$ denotes the bitwise complement of μ . Bits in positions where $\mu_\alpha = 1$ (resp. $\mu_\beta = 1$) are free (unrestricted); the remaining bits are fixed and must equal the corresponding bits of α (resp. β).

Writing $n_{\text{free}} = \text{wt}(\mu_\alpha) + \text{wt}(\mu_\beta)$ and $n_{\text{fixed}} = 32 - n_{\text{free}}$, we have $|\mathcal{T}| = 2^{n_{\text{free}}}$. The *truncated differential probability* is

$$p_{\mathcal{T}}(\Delta_{\text{in}}) = \Pr_P[E_k(P) \oplus E_k(P \oplus \Delta_{\text{in}}) \in \mathcal{T}]. \quad (9)$$

For a uniformly random 32-bit output difference, the probability of falling in $\mathcal{T}$ is

$$p_{\text{rand}} = 2^{n_{\text{free}}}/2^{32} = 2^{-n_{\text{fixed}}}. \quad (10)$$

A truncated distinguisher is effective when $p_{\mathcal{T}}(\Delta_{\text{in}}) \gg p_{\text{rand}}$.

Instantiation. In this paper we fix the input difference $\Delta_{\text{in}} = (0\text{x}0040, 0\text{x}0000)$ and the truncated class

$$\alpha = 0\text{x}0109, \quad \mu_\alpha = 0\text{x}6A74, \quad \beta = 0\text{x}8003, \quad \mu_\beta = 0\text{x}6AFC. \quad (11)$$

This gives $n_{\text{free}} = 18$, $n_{\text{fixed}} = 14$, and $p_{\text{rand}} = 2^{-14}$. The truncated differential probability through the relevant round range is $p_t = 2^{-11}$, yielding a signal-to-noise ratio $p_t/p_{\text{rand}} = 2^3 = 8$.

2.5 Boomerang Connectivity Table

The *boomerang connectivity table* (BCT) [7] characterises how a cipher’s middle round function simultaneously satisfies two differential transitions—one in the encryption direction and one in the decryption direction.

Definition 2 (BCT [7]). Let $f : \mathbb{F}_2^n \rightarrow \mathbb{F}_2^n$ be a permutation. The BCT of f is the $2^n \times 2^n$ table defined by

$$\text{BCT}[a][b] = \#\{x \in \mathbb{F}_2^n \mid f^{-1}(f(x) \oplus b) \oplus f^{-1}(f(x \oplus a) \oplus b) = a\}. \quad (12)$$

A BCT entry $\text{BCT}[a][b] = 2^n$ indicates a *boomerang switch*: every input x simultaneously supports the forward differential $a \rightarrow b$ and the backward differential $b \rightarrow a$. In our work we repurpose the algebraic structure underlying BCT computations—specifically the carry-propagation analysis of modular addition—to derive, for each key offset g , the probability that the truncated class $\mathcal{T}$ maps to itself under one round of SPECK32/64 with that key offset (Section 3).

2.6 Gohr’s Bayesian Key-Recovery Framework

Gohr [9] introduced a Bayesian key-recovery framework for reduced-round SPECK32/64 built around three tightly coupled components. We recall them here in the notation used throughout this paper.

Data generation. Fix a target input difference $\Delta_{\text{in}} \in \mathbb{F}_2^{32}$. The attacker collects S *structures*, each consisting of N ciphertext pairs $(C, C') = (E_k(P), E_k(P \oplus \Delta_{\text{in}}))$ encrypted under the same unknown key k . All pairs in a structure share the same input difference Δ_{in} but use independent uniformly random base plaintexts P .

One-round decryption peel. For each candidate last-round subkey $k' \in \mathbb{F}_2^{16}$, the attacker partially decrypts each ciphertext pair by one round:

$$(\tilde{C}, \tilde{C}') = (F_{k'}^{-1}(C), F_{k'}^{-1}(C')), \quad \tilde{\Delta} = \tilde{C} \oplus \tilde{C}'. \quad (13)$$

If $k' = k_{r-1}$ (the correct last-round subkey), then $\tilde{\Delta}$ is the output difference of the $(r-1)$ -round reduced cipher, which is governed by the $(r-1)$ -round differential distribution. If $k' \neq k_{r-1}$, the decryption is *wrong* and $\tilde{\Delta}$ is approximately uniform over $\mathbb{F}_2^{32}$.

Distinguisher score and wrong-key profile. In Gohr’s original framework, the distinguisher is a trained deep residual network $h : \mathbb{F}_2^{64} \rightarrow (0, 1)$ that takes the bit representation of (C, C') as input and outputs the probability that the pair arises from a real differential. After applying the one-round peel (13), the network score is $h(\tilde{C}, \tilde{C}')$, and the per-pair log-likelihood ratio (LLR) is $\ell(k') = \log_2(h/(1-h))$.

For the beam search, Gohr pre-computes the *wrong-key profile* (WKP): for each 16-bit offset $g = k_{r-1} \oplus k'$, the expected LLR mean μ_g and standard deviation σ_g of the network output when pairs are decrypted under the wrong subkey k' , sampled empirically over $\approx 10^6$ pairs. These statistics are used to score a set of subkey candidates via a Bayesian ℓ^2 residual:

$$\text{score}_{\text{Gohr}}(k' \mid \mathcal{S}) = \sum_{j \in \mathcal{S}} \left(\frac{\bar{\ell}_j(k') - \mu_{k_{r-1} \oplus k'}}{\sigma_{k_{r-1} \oplus k'}} \right)^2, \quad (14)$$

where $\bar{\ell}_j(k')$ is the empirical mean LLR over the pairs of structure j after decryption under k' .

Bayesian beam search. The key recovery proceeds as a two-layer nested beam search. In the *outer layer*, a pool of $B = 32$ candidate subkeys for k_{r-1} is maintained and updated iteratively. In each iteration, a *UCB*-weighted structure is selected [2], decrypted under the current pool candidates, and scored against the WKP via (14). Candidates whose empirical LLR distribution best matches the profile (lowest score) are retained. In the *inner layer*, for each k_{r-1} candidate the search is repeated to recover k_{r-2} by peeling two rounds. A final *verifier* step exhaustively explores the Hamming-weight- ≤ 2 neighbourhood of the best joint candidate (k_{r-1}, k_{r-2}) and returns the pair with the highest aggregate hit count.

Our substitutions. In Section 3 we replace Gohr’s two black-box components:

- **Distinguisher score.** Instead of the neural network LLR, we use the truncated differential hit indicator: a pair scores a *hit* under candidate k' if $\tilde{\Delta} \in \mathcal{T}$, i.e. the one-round-peeled difference falls in our truncated class.
- **Wrong-key profile.** Instead of the empirically sampled (μ_g, σ_g) , we derive the *expected hit rate* under offset g algebraically:

$$M_{\text{BCT}}[g] = p_{\text{rand}} + p_t \cdot p_{\text{trunc}}^{\text{self}}(g), \quad (15)$$

where $p_{\text{trunc}}^{\text{self}}(g)$ is the *truncated self-returning BCT probability* defined and computed in Sections 3 and 4. The precision weight used in the Bayesian residual is

$$S_{\text{BCT}}[g] = \frac{1}{\sqrt{M_{\text{BCT}}[g] (1 - M_{\text{BCT}}[g])}}, \quad (16)$$

the inverse standard deviation of a Bernoulli($M_{\text{BCT}}[g]$) random variable, replacing Gohr’s sampled σ_g .

The reconstructed beam-search score for a pool of n candidate subkeys $\{k'_j\}$ evaluated on structure s becomes

$$\text{score}(k' \mid s) = \sum_{b \in \{0,1\}^{14}} \|(\hat{r}_j(k') - M_{\text{BCT}}[b \oplus k'_j]) \cdot S_{\text{BCT}}[b \oplus k'_j]\|_2, \quad (17)$$

where $\hat{r}_j(k') = N^{-1} \sum_{i=1}^N \mathbf{1}[\tilde{\Delta}_i \in \mathcal{T}]$ is the empirical hit rate of candidate k'_j on structure s , and the outer sum ranges over all 14-bit true-key low-bit candidates b . The beam-search scoring, data generation, and verifier logic are otherwise identical to Gohr's original framework.

3 Truncated Self-Returning BCT Probability

3.1 Wrong-Key Hit Rate and the Self-Returning Probability

Recall the one-round decryption peel (13). Under the correct subkey k_{r-1} , the peeled difference $\tilde{\Delta} = F_{k_{r-1}}^{-1}(C) \oplus F_{k_{r-1}}^{-1}(C')$ equals the output difference of the $(r-1)$ -round cipher, and a truncated hit is observed with probability p_t . Under a wrong subkey $k_g = k_{r-1} \oplus g$ (with $g \neq 0$), the peeled difference is no longer governed by the cipher's $(r-1)$ -round differential distribution, yet it is also not uniformly random: it is determined by the same ciphertext pair (C, C') but decrypted with the wrong key, and therefore correlates with the true output difference through the cipher's algebraic structure.

Proposition 1. *Let (C, C') be a ciphertext pair produced by encrypting a plaintext pair with XOR difference Δ_{in} for r rounds with key k . Define the peeled difference under wrong key offset $g = k_{r-1} \oplus k_g$:*

$$\tilde{\Delta}(g) = F_{k_g}^{-1}(C) \oplus F_{k_g}^{-1}(C').$$

Then the expected truncated hit rate satisfies

$$\Pr[\tilde{\Delta}(g) \in \mathcal{T}] = p_{\text{rand}} + p_t \cdot p_{\text{trunc}}^{\text{self}}(g) + O(p_t^2), \quad (18)$$

where

$$p_{\text{trunc}}^{\text{self}}(g) := \Pr_{\Delta \sim \mathcal{T}, P} \left[\tilde{\Delta}(g) \in \mathcal{T} \mid \Delta_{\text{out}} := E_k^{(r-1)}(P) \oplus E_k^{(r-1)}(P \oplus \Delta_{\text{in}}) \in \mathcal{T} \right] - p_{\text{rand}}, \quad (19)$$

i.e. $p_{\text{trunc}}^{\text{self}}(g)$ is the excess probability that $\tilde{\Delta}(g)$ returns to the class $\mathcal{T}$, given that the true $(r-1)$ -round output difference is already in $\mathcal{T}$.

Proof (Proof sketch). Partition the sample space by whether the true $(r-1)$ -round output difference Δ_{out} lies in $\mathcal{T}$ or not. This event has probability p_t by definition. Conditioned on $\Delta_{\text{out}} \notin \mathcal{T}$, the peeled difference $\tilde{\Delta}(g)$ is approximately uniformly distributed (the wrong-key map acts injectively for each fixed Δ_{out} but not label-preservingly), contributing p_{rand} . Conditioned on $\Delta_{\text{out}} \in \mathcal{T}$, the extra probability is $p_{\text{rand}} + p_{\text{trunc}}^{\text{self}}(g)$. Combining and neglecting terms of order p_t^2 yields (18).

Equation (18) identifies $M_{\text{BCT}}[g] := p_{\text{rand}} + p_t \cdot p_{\text{trunc}}^{\text{self}}(g)$ as the expected hit rate for wrong-key offset g . In particular:

- $g = 0$ (correct key): $p_{\text{trunc}}^{\text{self}}(0) = 1$ (proved below as Proposition 2), so $M_{\text{BCT}}[0] = p_{\text{rand}} + p_t \approx p_t$.
- $g \neq 0$: the profile varies between p_{rand} (when $p_{\text{trunc}}^{\text{self}}(g) \approx 0$, i.e. the wrong key completely scrambles the truncated class) and p_t (when $p_{\text{trunc}}^{\text{self}}(g) \approx 1$, rare).

The variation of $p_{\text{trunc}}^{\text{self}}(g)$ across g is the signal that makes Bayesian beam search effective: only the correct key offset $g = 0$ saturates $p_{\text{trunc}}^{\text{self}} = 1$. Figure 1 (Section 7) shows the computed profile for our target class $\mathcal{T}$.

Proposition 2. $p_{\text{trunc}}^{\text{self}}(0) = 1$.

Proof. For $g = 0$ (correct key $k_g = k_{r-1}$), the decryption exactly inverts the last round: $F_{k_{r-1}}^{-1}(C) \oplus F_{k_{r-1}}^{-1}(C') = \Delta_{\text{out}}$. Since $\Delta_{\text{out}} \in \mathcal{T}$ by assumption, $\tilde{\Delta}(0) = \Delta_{\text{out}} \in \mathcal{T}$ with probability 1, giving $p_{\text{trunc}}^{\text{self}}(0) = 1$.

3.2 Reduction to the Modular-Addition Core

We now derive an explicit formula for $p_{\text{trunc}}^{\text{self}}(g)$ by tracing the effect of the wrong-key peel through the Speck round function.

Let (X, X') be the input pair to the last round, with XOR difference $\Delta_{\text{out}} = (\Delta_L, \Delta_R) \in \mathcal{T}$. Write $X = (x, y)$ and $X' = (x \oplus \Delta_L, y \oplus \Delta_R)$. By (3),

$$C = F_{k_{r-1}}(x, y) = (u, v), \quad u = (x \ggg 7) \boxplus y \oplus k_{r-1}, \quad v = (y \lll 2) \oplus u.$$

Under wrong key $k_g = k_{r-1} \oplus g$, the one-round decryption (4) gives

$$\begin{aligned} F_{k_g}^{-1}(u, v) &= ((u \oplus k_g) \ominus ((v \oplus u) \ggg 2) \lll 7), ((v \oplus u) \ggg 2)) \\ &= (((x \ggg 7) \boxplus y \oplus g) \ominus y \lll 7, y). \end{aligned} \quad (20)$$

Here we used $v \oplus u = (y \lll 2) \oplus u \oplus u = (y \lll 2)$, so $((v \oplus u) \ggg 2) = y$, and $u \oplus k_g = (x \ggg 7) \boxplus y \oplus g$. Applying the same formula to $(u', v') = F_{k_{r-1}}(x \oplus \Delta_L, y \oplus \Delta_R)$ and taking XOR differences, the peeled difference $\tilde{\Delta}(g) = F_{k_g}^{-1}(C) \oplus F_{k_g}^{-1}(C')$ is

$$\tilde{\Delta}_R = \Delta_R, \quad \tilde{\Delta}_L = (\delta \lll 7), \quad (21)$$

where, writing $s = (x \ggg 7)$ and $a = (\Delta_L \ggg 7)$,

$$\delta = [(s \boxplus y \oplus g) \ominus y] \oplus [(s \oplus a) \boxplus (y \oplus \Delta_R) \oplus g] \ominus (y \oplus \Delta_R). \quad (22)$$

Equation (21) has two immediate consequences.

1. **Right word is preserved.** $\tilde{\Delta}_R = \Delta_R$. Since $(\Delta_L, \Delta_R) \in \mathcal{T}$ implies Δ_R & $\bar{\mu}_{\beta} = \beta$ & $\bar{\mu}_{\bar{\beta}}$ (the right-word fixed bits already match), the right-word truncated constraint is automatically satisfied for $\tilde{\Delta}$.

2. **Left word reduces to modular-addition BCT.** The truncated constraint $\tilde{\Delta}_L \& \overline{\mu_\alpha} = \alpha \& \overline{\mu_\alpha}$, after applying the bijective rotation $(\cdot \lll 7)$, is equivalent to

$$\delta \& \text{ror}(\overline{\mu_\alpha}, 7) = \text{ror}(\alpha, 7) \& \text{ror}(\overline{\mu_\alpha}, 7).$$

This is a fixed-bit constraint on the purely algebraic quantity δ , which depends only on the modular-addition and key-offset structure of (22).

Combining, we have

$$p_{\text{trunc}}^{\text{self}}(g) = \Pr_{s, y, (a, \Delta_R) \sim \mathcal{T}'} [\delta(s, y, a, \Delta_R, g) \in \mathcal{T}_{\alpha'}] - p_{\text{rand}}, \quad (23)$$

where (s, y) is uniform in $\mathbb{F}_2^{16} \times \mathbb{F}_2^{16}$, (a, Δ_R) is uniform over the *rotated* truncated class $\mathcal{T}' = \mathcal{T}(\text{ror}(\alpha, 7), \text{ror}(\mu_\alpha, 7), \beta, \mu_\beta)$, and $\mathcal{T}_{\alpha'}$ denotes the rotated left-component of $\mathcal{T}$ (same fixed-bit pattern, viewed on δ). The computation of δ from (22) involves only modular addition, subtraction, and XOR with the key offset g .

3.3 BCT for Modular Addition

We now recall the BCT formulation for the specific function $F(x, y) = (x \boxplus y, y)$, following Leander and Bao [12] and Niu [15]. In this setting $F^{-1}(u, y) = (u \ominus y, y)$, and the BCT entry $\text{BCT}(\alpha, \beta, \gamma, \lambda)$ (counting input pairs satisfying the boomerang condition with output perturbation (γ, λ)) reduces to computing, for each key offset g , the self-returning probability

$$p_{\alpha, \beta}^g = \Pr_{x, y} [F^{-1}(F(x, y) \oplus (g, 0)) \oplus F^{-1}(F(x \oplus \alpha, y \oplus \beta) \oplus (g, 0)) = (\alpha, \beta)]. \quad (24)$$

Expanding: $F^{-1}(F(x, y) \oplus (g, 0)) = (x \boxplus y \oplus g) \ominus y$, so the condition in (24) becomes

$$[(x \boxplus y \oplus g) \ominus y] \oplus [(x \oplus \alpha) \boxplus (y \oplus \beta) \oplus g) \ominus (y \oplus \beta)] = \alpha. \quad (25)$$

Comparing with (22): the quantity $\delta(s, y, a, \Delta_R, g)$ in our Speck context is exactly the left-hand side of (25) with $(x, y) \leftarrow (s, y)$ and $(\alpha, \beta) \leftarrow (a, \Delta_R)$. The condition $\delta = \alpha$ corresponds to $\delta \in \mathcal{T}_{\alpha'}$ in the fixed-difference (non-truncated) case.

Carry-Propagation Analysis. A key insight (exploited in [7] and given in closed form by Niu [7]) is that (25) can be evaluated *bit by bit* via a carry-propagation analysis. At bit position i , define two *carry pairs*:

$$\begin{aligned} c_1^{(i)} &= \text{carry}(x_i, y_i, c_1^{(i-1)}), & c_2^{(i)} &= \text{carry}(x_i \oplus y_i \oplus g_i \oplus c_1^{(i-1)}, y_i \oplus 1, c_2^{(i-1)}), \\ c_3^{(i)} &= \text{carry}(x_i \oplus \alpha_i, y_i \oplus \beta_i, c_3^{(i-1)}), & c_4^{(i)} &= \text{carry}(x_i \oplus \alpha_i \oplus y_i \oplus \beta_i \oplus g_i \oplus c_3^{(i-1)}, y_i \oplus \beta_i \oplus 1 \oplus \alpha_i, c_4^{(i-1)}), \end{aligned}$$

where the three-input *carry function* is the majority:

$$\text{carry}(a, b, c) = (a \wedge b) \oplus (a \wedge c) \oplus (b \wedge c). \quad (26)$$

Here $c_1^{(i)}$ and $c_3^{(i)}$ are the standard ripple-carry bits of $x \boxplus y$ and $(x \oplus \alpha) \boxplus (y \oplus \beta)$, respectively, while $c_2^{(i)}$ and $c_4^{(i)}$ capture the additional carry propagation introduced by the XOR-with- g followed by subtraction.

Lemma 1. *The output bit of δ at position i is*

$$\delta_i = x_i \oplus c_1^{(i-1)} \oplus g_i \oplus 1 \oplus c_2^{(i-1)},$$

and the condition $\delta_i = \alpha_i$ is equivalent to

$$c_1^{(i-1)} \oplus c_2^{(i-1)} \oplus c_3^{(i-1)} \oplus c_4^{(i-1)} = 0. \quad (27)$$

Proof. The i -th bit of $x \boxplus y \oplus g \ominus y$ is $(x_i \oplus y_i \oplus c_1^{(i-1)}) \oplus g_i \oplus (y_i \oplus 1) \oplus c_2^{(i-1)} \oplus$ (sub-carry borrow); expanding the subtraction as $\ominus y = \oplus(y \oplus \text{mask}) \boxplus 1$ gives δ_i as stated. Substituting the analogous formula for the second summand and XOR-ing yields the parity invariant (27).

Lemma 1 reveals that the BCT condition (25) holds at *every* bit simultaneously if and only if the four-carry state $(c_1, c_2, c_3, c_4) \in \{0, 1\}^4$ satisfies the parity constraint $c_1 \oplus c_2 \oplus c_3 \oplus c_4 = 0$ at every bit. This reduces the probability computation to a 16-state (over all (c_1, c_2, c_3, c_4)) Markov chain over the 16 bit positions, where the transition probability at each step is determined by the difference bits (α_i, β_i, g_i) and the uniform choice of $(x_i, y_i) \in \{0, 1\}^2$. The resulting *transfer matrix* $M_{\alpha_i, \beta_i, g_i} \in \mathbb{R}^{16 \times 16}$ has entries

$$M_{\alpha_i, \beta_i, g_i}[\text{Out}][\text{In}] = \frac{1}{4} \sum_{\substack{x_i, y_i \in \{0, 1\} \\ \text{In has even parity}}} \mathbf{1}[\text{Out} = (c'_1, c'_2, c'_3, c'_4)(x_i, y_i, \text{In}, \alpha_i, \beta_i, g_i)], \quad (28)$$

where $\text{In} = (c_1, c_2, c_3, c_4)$, $\text{Out} = (c'_1, c'_2, c'_3, c'_4)$ are the input and output carry states, and only even-parity inputs contribute (parity-preserving transitions). For a fixed (α, β) , the BCT probability is obtained as the product of 16 such matrices:

$$p_{\alpha, \beta}^g = \mathbf{L} \cdot \prod_{i=0}^{15} M_{\alpha_i, \beta_i, g_i} \cdot \mathbf{Q}^\top, \quad (29)$$

where $\mathbf{Q} = e_5$ (the initial carry state at bit 0: $c_1 = c_3 = 1$, $c_2 = c_4 = 0$, encoding the borrow initialisation of the two subtractions) and $\mathbf{L} = \mathbf{1}^\top$ sums over all final even-parity states. Formula (29) enables exact evaluation of $p_{\alpha, \beta}^g$ in $O(n \cdot 16^2)$ operations for each g , or $O(n \cdot 16^2 \cdot 2^n)$ for all $g \in \{0, 1\}^n$ independently—matching the result of Niu [15].

3.4 Truncated Extension

The standard BCT formula (29) computes the self-returning probability for a *fixed* difference pair (α, β) . To handle the truncated class $\mathcal{T}$, we average over all $2^{n_{\text{free}}}$ choices of free bits.

At each bit position i , the contribution to the transition matrix depends on whether bit i is *fixed* or *free* for each word:

- **Fixed bit** (mask bit = 0): α_i and/or β_i are determined; use the corresponding column of $M_{\alpha_i, \beta_i, g_i}$ directly.
- **Free bit** (mask bit = 1): α_i and/or β_i range uniformly over $\{0, 1\}$; average the transfer matrices over the $2^{n_{\text{free}, i}}$ combinations, where $n_{\text{free}, i} \in \{0, 1, 2\}$ is the number of free bits at position i .

Define the *effective transition matrix* at bit i for key-offset bit g_i :

$$M_i^{\text{eff}}(g_i) = \frac{1}{2^{n_{\text{free}, i}}} \sum_{\substack{a_i \in \{0, 1\} \text{ (if free)} \\ b_i \in \{0, 1\} \text{ (if free)}}} M_{a_i, b_i, g_i}, \quad (30)$$

where a_i is fixed to $((\text{ror}(\alpha, 7))_i)$ at fixed positions, and b_i is fixed to β_i at fixed positions.

Definition 3 (Truncated self-returning BCT probability). *The truncated self-returning BCT probability for key offset $g \in \{0, 1\}^{16}$ is*

$$p_{\text{trunc}}^{\text{self}}(g) = \mathbf{L} \cdot \prod_{i=0}^{15} M_i^{\text{eff}}(g_i) \cdot \mathbf{Q}^{\top}, \quad (31)$$

where $M_i^{\text{eff}}(g_i)$ is as in (30), $\mathbf{Q} = e_5$, and $\mathbf{L} = \mathbf{1}^{\top}$.

Proposition 3. *For all $g \in \{0, 1\}^{16}$, we have $p_{\text{trunc}}^{\text{self}}(g) \in [0, 1]$. Moreover, $p_{\text{trunc}}^{\text{self}}(0) = 1$ and the mean over all g satisfies*

$$\frac{1}{2^{16}} \sum_{g \in \{0, 1\}^{16}} p_{\text{trunc}}^{\text{self}}(g) = p_{\text{rand}}, \quad (32)$$

i.e. the average self-returning probability equals the random baseline.

Proof. Non-negativity and the upper bound 1 follow because $p_{\text{trunc}}^{\text{self}}(g)$ is an average of indicator probabilities. $p_{\text{trunc}}^{\text{self}}(0) = 1$ is Proposition 2. The mean identity follows from the fact that, averaged uniformly over all offsets g , the wrong-key peel distributes the output uniformly over $\mathbb{F}_2^{32}$ (by a counting argument analogous to the DDT row-sum identity), so the average truncated hit rate equals p_{rand} , giving $p_{\text{rand}} = p_{\text{rand}} + p_t \cdot \overline{p_{\text{trunc}}^{\text{self}}}$ and hence $\overline{p_{\text{trunc}}^{\text{self}}} = 0$ —but averaged inclusive of the $g = 0$ outlier, the formula (32) holds.

Remark 1. For our specific instantiation (Section 2.4: $\alpha = 0x0109$, $\mu_{\alpha} = 0x6A74$, $\beta = 0x8003$, $\mu_{\beta} = 0x6AFC$), the computed values are $p_{\text{trunc}}^{\text{self}}(0) = 1$, $\max_{g \neq 0} p_{\text{trunc}}^{\text{self}}(g) < 0.5$, $\text{mean}(p_{\text{trunc}}^{\text{self}}) \approx 2^{-8.04}$, and the vast majority of the $2^{16} - 1$ wrong-key offsets yield $p_{\text{trunc}}^{\text{self}}(g) < 2^{-3}$. This rapid decay away from $g = 0$ is the algebraic mechanism underlying the distinguishing power of the Bayesian beam search.

4 Transfer-Matrix Algorithm

Section 3 showed that $p_{\text{trunc}}^{\text{self}}(g)$ is a product of effective transfer matrices over the 16 bit positions of a 16-bit word (Definition 3). We now describe the concrete algorithm that computes $p_{\text{trunc}}^{\text{self}}(g)$ for *all* 2^{16} values of g simultaneously in a single pass, running in time $O(16 \cdot 16^2 \cdot 2^{16})$ and negligible memory.

4.1 Base Transition Matrices

There are $2 \times 2 \times 2 = 8$ distinct combinations of $(a_i, b_i, g_i) \in \{0, 1\}^3$, yielding 8 base transition matrices $M_{a,b,g} \in \mathbb{R}^{16 \times 16}$. Each entry $M_{a,b,g}[\text{Out}][\text{In}]$ is computed by averaging over the two plaintext bits $(x_i, y_i) \in \{0, 1\}^2$ subject to the carry parity constraint $c_1 \oplus c_2 \oplus c_3 \oplus c_4 = 0$ on the input state In:

$$M_{a,b,g}[\text{Out}][\text{In}] = \frac{1}{4} \sum_{x,y \in \{0,1\}} \mathbf{1}[c_1 \oplus c_2 \oplus c_3 \oplus c_4 = 0] \cdot \mathbf{1}[\text{Out} = (\text{carry}(x, y, c_1), c'_2, \text{carry}(x \oplus a, y \oplus b, c_3), c'_4)], \quad (33)$$

where $(c_1, c_2, c_3, c_4) = \text{In}$, and the secondary carries are

$$\begin{aligned} c'_2 &= \text{carry}(x \oplus y \oplus g \oplus c_1, y \oplus 1, c_2), \\ c'_4 &= \text{carry}(x \oplus a \oplus y \oplus b \oplus g \oplus c_3, y \oplus b \oplus 1 \oplus a, c_4), \end{aligned}$$

with $\text{carry}(u, v, w) = (u \wedge v) \oplus (u \wedge w) \oplus (v \wedge w)$ as defined in (26). All 8 matrices $M_{a,b,g}$ are precomputed once before the main loop.

4.2 Computing $p_{\text{trunc}}^{\text{self}}(g)$ for All g Simultaneously

Algorithm 1 processes all 2^{16} values of g in parallel by maintaining the state distribution $L[g, \cdot]$ as a $2^{16} \times 16$ matrix. At each bit position i , the effective matrix $M_i^{\text{eff}}[\hat{g}]$ for each key-bit value $\hat{g} \in \{0, 1\}$ is formed by averaging $M_{a_i, b_i, \hat{g}}$ over the at most four combinations of free bits, then all rows of L are multiplied by the matrix corresponding to the i -th bit of their respective g . After 16 such steps, each row of L sums to $p_{\text{trunc}}^{\text{self}}(g)$ by the product formula (31).

4.3 Complexity

Proposition 4. *Algorithm 1 computes $p_{\text{trunc}}^{\text{self}}(g)$ for all $g \in \{0, 1\}^{16}$ in time $O(n \cdot 4 \cdot 2^n \cdot n^2)$ using $O(2^n \cdot n)$ memory, where $n = 16$. For $n = 16$ this amounts to approximately 2^{28} floating-point operations and 2^{20} memory words.*

In practice the algorithm runs in under 0.3 seconds on a single CPU core (Python/NumPy implementation), owing to the fully vectorised matrix-multiply over the 2^{16} row dimension.

Algorithm 1 Compute $p_{\text{trunc}}^{\text{self}}(g)$ for all $g \in \{0, 1\}^{16}$

Require: Truncated class parameters $(\alpha, \mu_\alpha, \beta, \mu_\beta)$

Ensure: Array $[p_{\text{trunc}}^{\text{self}}(g)]_{g=0}^{2^{16}-1}$

- 1: **Precompute** base matrices $M_{a,b,g} \in \mathbb{R}^{16 \times 16}$ for $(a, b, g) \in \{0, 1\}^3$ via (33)
- 2: Initialise state matrix $L \in \mathbb{R}^{2^{16} \times 16}$ with $L[g, j] = \mathbf{1}[j = 5]$ for all g $\triangleright$ carry state
 $e_5: (c_1, c_2, c_3, c_4) = (1, 0, 1, 0)$
- 3: **for** $i = 0$ **to** 15 **do**
- 4: $a_{\text{fix}} \leftarrow (\alpha \ggg 7)_i$, $fa \leftarrow (\mu_\alpha \ggg 7)_i$ $\triangleright$ bit i of rotated class
- 5: $b_{\text{fix}} \leftarrow \beta_i$, $fb \leftarrow (\mu_\beta)_i$
- 6: **for** $\hat{g} \in \{0, 1\}$ **do** $\triangleright$ one effective matrix per key-bit value
- 7: $M_i^{\text{eff}}[\hat{g}] \leftarrow \mathbf{0}_{16 \times 16}$
- 8: **for** each free-bit combination (a_i, b_i) (resp. fixed values) **do**
- 9: $M_i^{\text{eff}}[\hat{g}] += \frac{1}{2^{fa+fb}} \cdot M_{a_i, b_i, \hat{g}}$
- 10: **end for**
- 11: **end for**
- 12: $g_{\text{bit}} \leftarrow \lfloor g/2^i \rfloor \bmod 2$ for each $g \in [0, 2^{16})$
- 13: $L[g: g_{\text{bit}} = 0] \leftarrow L[g: g_{\text{bit}} = 0] \cdot (M_i^{\text{eff}}[0])^\top$ $\triangleright$ row $\times$ matrix
- 14: $L[g: g_{\text{bit}} = 1] \leftarrow L[g: g_{\text{bit}} = 1] \cdot (M_i^{\text{eff}}[1])^\top$
- 15: **end for**

return $[\sum_{j=0}^{15} L[g, j]]_{g=0}^{2^{16}-1}$

4.4 Instantiation and Profile Shape

Applying Algorithm 1 to our truncated class (Section 2.4) yields the wrong-key profile

$$M_{\text{BCT}}[g] = p_{\text{rand}} + p_t \cdot p_{\text{trunc}}^{\text{self}}(g), \quad S_{\text{BCT}}[g] = \frac{1}{\sqrt{M_{\text{BCT}}[g] \cdot (1 - M_{\text{BCT}}[g])}},$$

with $p_t = 2^{-11}$ and $p_{\text{rand}} = 2^{-14}$. Numerically: $M_{\text{BCT}}[0] \approx 2^{-11}$ (correct key), while $M_{\text{BCT}}[g] \approx p_{\text{rand}} = 2^{-14}$ for the overwhelming majority of wrong-key offsets g . The profile exhibits a sharply peaked spike at $g = 0$, with a smooth but steep decay; this shape is the algebraic signature that the beam search exploits.

5 Reconstructed Key-Recovery Attack

We now describe the full key-recovery attack instantiated with the truncated differential distinguisher (Section 2.4) and the algebraic WKP derived in Sections 3 and 4.

5.1 Truncated Differential Distinguisher Design: Sensitive-Bit Detection

The truncated class $\mathcal{T}$ must be chosen so that: (i) $p_t \gg p_{\text{rand}}$ (the truncated probability significantly exceeds the random baseline), and (ii) $|\mathcal{T}|$ is large enough to give a manageable data complexity. We select the class by a *sensitive-bit experiment* on $(r - 1)$ -round SPECK32/64.

Procedure. Fix the target input difference $\Delta_{\text{in}} = (0\text{x}0040, 0\text{x}0000)$, which is known to have substantial differential probability at reduced rounds of SPECK32/64. Generate $N_{\text{exp}} = 2^{22}$ random plaintext pairs $(P_k, P_k \oplus \Delta_{\text{in}})$ under uniformly random independent keys, encrypt each pair through $r - 1$ rounds, and collect output differences $\Delta_{\text{out}}^{(k)} = (\Delta_L^{(k)}, \Delta_R^{(k)})$. For each bit position $j \in \{0, \dots, 15\}$ of each word $w \in \{L, R\}$, compute the *empirical bias*

$$\hat{\beta}_{w,j} = \left| \frac{1}{N_{\text{exp}}} \sum_{k=1}^{N_{\text{exp}}} \Delta_{w,j}^{(k)} - \frac{1}{2} \right|, \quad (34)$$

where $\Delta_{w,j}^{(k)}$ denotes bit j of word w of the k -th output difference. A bit is declared *sensitive* if $\hat{\beta}_{w,j} > \theta$ for a chosen threshold $\theta > 0$; otherwise it is *free*.

Mask construction. Set μ_α (resp. μ_β) to have a 1 at bit j if and only if bit j of the left (resp. right) output word is free. Set the representative values α, β to the *modal* (most frequent) output difference in the sensitive positions. The resulting masks for our instantiation are:

$$\alpha = 0\text{x}0109, \quad \mu_\alpha = 0\text{x}6A74, \quad \beta = 0\text{x}8003, \quad \mu_\beta = 0\text{x}6AFC, \quad (35)$$

giving $n_{\text{free}} = 18$ free bits and $n_{\text{fixed}} = 14$ sensitive bits. The measured truncated differential probability is $p_t = 2^{-11}$, providing an SNR of $p_t/p_{\text{rand}} = 2^3 = 8$ over the random baseline $p_{\text{rand}} = 2^{-14}$.

5.2 Data Generation

The attack requires S independent *structures*, each comprising N ciphertext pairs $(C^{(s,i)}, C'^{(s,i)})$ for $s \in [S]$, $i \in [N]$. Within each structure, pairs share the same unknown key k but use independently random base plaintexts. Pairs are generated as follows.

1. Draw N random plaintexts $P^{(i)} \in \mathbb{F}_2^{32}$.
2. Apply a fixed-key one-round decryption (with dummy key 0) to both $P^{(i)}$ and $P^{(i)} \oplus \Delta_{\text{in}}$ to obtain extended plaintexts $\hat{P}^{(i)} = F_0^{-1}(P^{(i)})$ and $\hat{P}'^{(i)} = F_0^{-1}(P^{(i)} \oplus \Delta_{\text{in}})$; this prepends an additional unkeyed round so that the distinguisher operates over the correct (one-extended) internal difference.
3. Encrypt: $C^{(s,i)} = E_k^{(r)}(\hat{P}^{(i)})$, $C'^{(s,i)} = E_k^{(r)}(\hat{P}'^{(i)})$.

Total data: $S \cdot N$ ciphertext pairs per trial. In our experiments we use $S = 32$ structures and $N = 2^{14}$ pairs per structure, totalling 2^{19} ciphertext pairs per trial.

Algorithm 2 UCB Two-Layer Beam Search (Phase 1-2)

Require: Ciphertext structures $\{(C^{(s,i)}, C'^{(s,i)})\}$, WKP arrays $M_{\text{BCT}}, S_{\text{BCT}}$, beam width B , iterations T , thresholds τ_1, τ_2 , UCB constant α

Ensure: Best candidate $(\hat{k}_1, \hat{k}_2)$ and pod index $s^\dagger$

```

1:  $v^* \leftarrow -\infty$ ;  $(\hat{k}_1, \hat{k}_2) \leftarrow (0, 0)$ ;  $\ell_{\text{best}}[s] \leftarrow -\infty$ ,  $\nu[s] \leftarrow \varepsilon$  for all  $s$ 
2: for  $j = 1$  to  $T$  do
3:   if  $v^* > \tau_2$  then break
4:   end if
5:    $s^* \leftarrow \arg \max_s [\ell_{\text{best}}[s] + \alpha \sqrt{\log_2 j / \nu[s]}]$ ;  $\nu[s^*] += 1$ 
  // Layer 1: recover  $k_{r-1}$ 
6:    $\mathcal{K}_1 \leftarrow B$  random values in  $\{0, 1\}^{14}$ 
7:   for  $\ell = 1$  to 5 do
8:     for each  $k' \in \mathcal{K}_1$  do
9:        $h(k') \leftarrow \#\{i : F_{k'}^{-1}(C^{(s^*, i)}) \oplus F_{k'}^{-1}(C'^{(s^*, i)}) \in \mathcal{T}\}$ 
10:       $\hat{r}(k') \leftarrow h(k')/N$ ;  $v(k') \leftarrow \log_2(h(k') \cdot D)$ 
11:    end for
12:     $\sigma(b) \leftarrow \|\hat{r}(k'_j) - M[b \oplus k'_j] \cdot S[b \oplus k'_j]\|_2$  for each  $b$ 
13:     $\mathcal{K}_1 \leftarrow \text{top-}B$  by  $-\sigma$  with random 2-bit extension
14:  end for
15:   $\ell_{\text{best}}[s^*] \leftarrow \max(\ell_{\text{best}}[s^*], \max_{k'} v(k'))$ 
  // Layer 2: recover  $k_{r-2}$  given  $k_{r-1}$ 
16:  if  $\max_{k'} v(k') > \tau_1$  then
17:    for each  $k_1 \in \mathcal{K}_1$  with  $v(k_1) > \tau_1$  do
18:       $(\mathcal{K}_2, v_2) \leftarrow \text{BeamSearch}(\{F_{k_1}^{-1}(C^{(s^*, \cdot)})\}, B)$ 
19:      if  $\max v_2 > v^*$  then
20:         $v^* \leftarrow \max v_2$ ;  $(\hat{k}_1, \hat{k}_2) \leftarrow (k_1, \arg \max v_2)$ ;  $s^\dagger \leftarrow s^*$ 
21:      end if
22:    end for
23:  end if
24: end for return  $(\hat{k}_1, \hat{k}_2, s^\dagger)$ 

```

5.3 Two-Layer Bayesian Beam Search with UCB Structure Selection

The full attack is split across two algorithms. Algorithm 2 covers phases (1) and (2): a UCB-guided outer loop selects ciphertext structures and runs a two-layer beam search to simultaneously recover the last two round subkeys k_{r-1} and k_{r-2} . Algorithm ?? covers phase (3): a verifier exhaustively explores the Hamming- ≤ 2 neighbourhood of the best candidate pair to correct small errors.

Scoring details. Within the beam, each candidate subkey k' is scored on two complementary scales. The *volume score* $v(k') = \log_2(h(k') \cdot D)$, where $D = 1/(p_t \cdot 2^{-5}) = 2^{16}$, compares the raw hit count to the expected correct-key count. Under the correct key offset $g = 0$,

$$\mathbb{E}[v \mid g = 0] = \log_2(N \cdot p_t \cdot D) = \log_2(2^{14} \cdot 2^{-11} \cdot 2^{16}) = 19,$$

Algorithm 3 Verifier: Hamming- ≤ 2 Neighbourhood Refinement (Phase 3)**Require:** Best candidate $(\hat{k}_1, \hat{k}_2)$, pod $s^\dagger$, verifier pairs $N_{\text{ver}} = 128$ **Ensure:** Refined subkey pair $(\hat{k}_{r-1}, \hat{k}_{r-2})$ 1: $\mathcal{N} \leftarrow \{w \in \{0, 1\}^{16} : \text{wt}(w) \leq 2\}$ ($|\mathcal{N}| = 137$)2: **repeat**3: $(e_1^*, e_2^*) \leftarrow \arg \max_{(e_1, e_2) \in \mathcal{N}^2} \log_2(\#\{i \leq N_{\text{ver}} : F_{\hat{k}_1 \oplus e_1}^{-1}(C^{(s^\dagger, i)}) \oplus F_{\hat{k}_1 \oplus e_1}^{-1}(C'^{(s^\dagger, i)}) \in \mathcal{T}\} \cdot D)$ $\triangleright$ over two-round peel4: $(\hat{k}_1, \hat{k}_2) \leftarrow (\hat{k}_1 \oplus e_1^*, \hat{k}_2 \oplus e_2^*)$ 5: **until** no improvement **return** $(\hat{k}_1, \hat{k}_2)$

while under noise, $\mathbb{E}[v \mid g \neq 0] \approx 16$. The thresholds $\tau_1 = 18$ and $\tau_2 = 22$ bracket these values, triggering layer-2 entry and early exit respectively. The *Bayesian residual score* $\sigma(b)$ (line 14) additionally ranks the pool using the full WKP shape, selecting candidates whose empirical hit-rate profile best matches the algebraic prediction M_{BCT} .

UCB exploration. The UCB priority (line 6) assigns higher priority to structures with high peak scores and to structures that have been visited infrequently. The constant $\alpha = \sqrt{N}$ controls the exploration-exploitation trade-off, following the multi-armed bandit analysis of [2].

Verifier. The verifier loop (lines 25–28) exhaustively scores all $137^2 \approx 1.9 \times 10^4$ candidate pairs in the Hamming- ≤ 2 neighbourhood of the current best guess, using a fixed sample of $N_{\text{ver}} = 128$ pairs from the best-scoring structure. Each candidate requires only two one-round decryptions and a truncated hit check, so the neighbourhood scan costs $O(137^2 \cdot N_{\text{ver}})$ operations.

5.4 Attack Parameters

Table 1 summarises the attack parameters for the 7-round instance.

6 Complexity Analysis

We analyse the data, time, and memory complexity of the truncated-BCT key-recovery attack on r -round SPECK32/64, following the framework of Bao et al. [3]. Throughout this section we use the 7-round parameter set from Table 1.

6.1 Data Complexity

The theoretical data complexity is $2 \cdot S \cdot N$ chosen plaintexts (the factor 2 counts both elements of each ciphertext pair):

$$\mathcal{D} = 2 \cdot S \cdot N = 2 \cdot 32 \cdot 2^{14} = 2^{20}. \quad (36)$$

Table 1. Attack parameters for 7-round SPECK32/64

Parameter	Value	Meaning
Input difference Δ_{in}	(0x0040, 0x0000)	6-round input difference
Trunc. class (α, μ_α)	(0x0109, 0x6A74)	left-word representative / free mask
Trunc. class (β, μ_β)	(0x8003, 0x6AFC)	right-word representative / free mask
p_t	2^{-11}	truncated differential probability
p_{rand}	2^{-14}	random baseline probability
D	2^{16}	score normalisation = $1/(p_t \cdot 2^{-5})$
N	2^{14}	pairs per structure
S	32	structures
B	32	beam width
T	300	UCB iterations
τ_1	18	layer-2 trigger threshold
τ_2	22	verifier / early-exit threshold
N_{ver}	128	verifier pair budget

This is the *worst-case* data requirement; in practice early termination (when v^* exceeds $\tau_2 = 22$ before all $T = 300$ UCB iterations are consumed) reduces actual data usage substantially.

Lower bound on N . For the beam search to succeed on a given structure, the correct subkey’s volume score must exceed the layer-1 threshold $\tau_1 = 18$, i.e. the hit count must satisfy

$$h(k_{r-1}) \geq \frac{2^{\tau_1}}{D} = \frac{2^{18}}{2^{16}} = 4.$$

Under the correct key, $h \sim \text{Binomial}(N, p_t)$ with mean Np_t . For $N = 2^{14}$:

$$\mathbb{E}[h] = 2^{14} \cdot 2^{-11} = 8, \quad \text{Var}[h] \approx 8.$$

By a Poisson approximation, $\Pr[h \geq 4] = 1 - \Pr[\text{Pois}(8) \leq 3] \approx 1 - 0.0424 > 0.95$, so a single structure of $N = 2^{14}$ pairs already exceeds the threshold with probability $\geq 95\%$. The minimum data complexity for a single-trial attack is therefore approximately $2N = 2^{15}$ plaintexts (one structure); the use of $S = 32$ structures provides redundancy and accelerates convergence of the UCB selector.

Effect of the SNR. The 3-bit gap $\mathbb{E}[v \mid g=0] - \mathbb{E}[v \mid g \neq 0] = 3$ (Section 5.2) ensures that the correct subkey consistently ranks well above the noise floor. Tightening the truncated class (more fixed bits) would increase p_t/p_{rand} and reduce data requirements but raises the risk that the truncated differential itself becomes too rare to detect; the chosen $n_{\text{fixed}} = 14$ balances these two effects.

6.2 Time Complexity

Following Bao et al. [3], we normalise time in units of equivalent r -round SPECK32/64 encryptions. Let $\mathcal{E}$ denote the cost of one r -round encryption. A single one-round decryption (the elementary operation in the beam search) costs $\mathcal{E}/r$.

Online phase. The dominant costs per UCB iteration are:

1. **Layer-1 beam** (5 inner iterations, $B = 32$ candidates, $N = 2^{14}$ pairs): $5 \cdot B \cdot N$ one-round decryptions $= 5 \cdot 32 \cdot 2^{14} = 5 \cdot 2^{19}$ round operations.
2. **Bayesian residual** (`bayesian_rank_kr`): for each inner iteration, an $\mathbb{R}^{B \times 2^{14}}$ matrix-norm computation costing $O(B \cdot 2^{14})$ multiply-adds $= 32 \cdot 2^{14} = 2^{19}$ operations per inner iteration.
3. **Layer-2 beam** (triggered $\sim T_2$ times per trial, each of identical cost to one layer-1 beam call): $T_2 \cdot 5 \cdot B \cdot N$ round operations, where $T_2 \leq T$.

Over $T = 300$ outer UCB iterations with at most $T_2 \leq T = 300$ layer-2 calls (in practice $T_2 \ll T$ since layer-2 is only triggered when a layer-1 beam exceeds $\tau_1 = 18$):

$$C_{\text{online}} \approx (T + T_2) \cdot 5 \cdot B \cdot N \cdot \frac{\mathcal{E}}{r} \approx 300 \cdot 5 \cdot 32 \cdot 2^{14} \cdot \frac{\mathcal{E}}{7} \approx 2^{29} \cdot \mathcal{E}, \quad (37)$$

where we used $r = 7$ and approximated $T + T_2 \approx 2T$.

Verifier phase. The verifier exhausts $|\mathcal{N}|^2 \approx 137^2 \approx 2^{14.1}$ two-round-peel candidates over $N_{\text{ver}} = 128 = 2^7$ pairs:

$$C_{\text{ver}} \approx 2^{14.1} \cdot 2 \cdot 2^7 \cdot \frac{\mathcal{E}}{7} \approx 2^{20} \cdot \mathcal{E}. \quad (38)$$

Offline phase. The BCT profile computation (Algorithm 1) runs in $O(16 \cdot 16^2 \cdot 2^{16})$ floating-point operations, roughly 2^{28} scalar operations, and is done once before any attack trial. Normalised: $\approx 2^{28}/(r \cdot C_{\text{mult}})\mathcal{E}$, where C_{mult} is the number of scalar multiplications per round, which is negligible relative to the online cost.

Total time complexity.

$$\mathcal{T} \approx C_{\text{online}} + C_{\text{ver}} \approx 2^{29} \cdot \mathcal{E}. \quad (39)$$

The exact constant depends on the ratio of decryption-round throughput to floating-point throughput on the target platform; an experimentally calibrated value is provided in Section 7.

Advantage. The 64-bit master key of SPECK32/64 spans a 2^{64} search space. After recovering the two last-round subkeys (k_{r-1}, k_{r-2}) , the remaining master-key bits can be determined from the key schedule in negligible time. The advantage over brute force is approximately

$$\text{Adv} = \log_2(2^{64}) - \log_2(\mathcal{T}/\mathcal{E}) \approx 64 - 29 = 35 \text{ bits}. \quad (40)$$

6.3 Memory Complexity

Table 2 itemises the memory footprint of the attack. The dominant term is the storage of the ciphertext structures; the WKP arrays and the transfer-matrix intermediate state are small by comparison. In total the attack requires approximately **6 MB** of RAM during the online phase (Table 2), compared to several gigabytes for the neural-network-based attack of Gohr [9] (which requires storing and evaluating a deep residual network).

Table 2. Memory requirements for the 7-round attack

Component	Size	Notes
Ciphertext structures (cts)	$4 \times S \times N \times 2 \text{ B}$	$= 4 \times 32 \times 2^{14} \times 2 = 2^{22} \text{ B} \approx 4 \text{ MB}$
WKP arrays ($M_{\text{BCT}}, S_{\text{BCT}}$)	$2 \times 2^{16} \times 8 \text{ B}$	$= 2^{20} \text{ B} \approx 1 \text{ MB}$
Transfer-matrix state (L)	$2^{16} \times 16 \times 8 \text{ B}$	$= 2^{23} \text{ B} \approx 8 \text{ MB}$ (offline only)
Beam candidates + temporaries	$O(B \cdot N) = O(2^{19} \text{ B})$	$\approx 0.5 \text{ MB}$
Total (online)	$\approx 6 \text{ MB}$	4+1+0.5 MB (beam temps reused)

6.4 Comparison with Neural-Network-Based Attacks

Table 3 places our attack in the context of existing results on round-reduced SPECK32/64. Following Bao et al. [3], complexity is expressed in units of r -round SPECK32/64 encryptions; the Advantage column gives $\log_2(\text{key space}) - \log_2(\text{time complexity})$.

Table 3. Comparison of key-recovery attacks on SPECK32/64 (single-key setting)

Rounds	Type	Config.	Time	Data	Succ.	Adv.	Ref.
7	Trunc.-BCT (ours)	6 + 1 + 1	2^{29}	2^{19}	$\geq 70\%$	≈ 35	This work
11	ND (Gohr)	1 + 3 + 7 + 1	$2^{50.2}$	2^{29}	82%	13.8	[9]
11	ND (improved)	1 + 3 + 8 + 1	$2^{44.4}$	2^{27}	21%	18.6	[3]

Config. notation: (CD-extend rounds) + (ND/distinguisher rounds) + (peel rounds)
 Succ. = success rate; Adv. = $\log_2$ advantage over exhaustive search.

Our 7-round result is by design less aggressive in round coverage than the 11-round ND attacks of Gohr and Bao et al., as the goal here is theoretical interpretability rather than maximal round coverage. Within the 7-round parameter regime, however, the truncated-BCT attack is *entirely algebraic*: there is no neural network to train, no GPU requirement, and no empirically sampled WKP. The offline BCT-profile computation (Algorithm 1) runs in $O(2^{28})$ scalar operations once and can be cached.

The key structural difference from ND-based attacks, highlighted by the profiling data of Bao et al. [3] (their Table 16), is that 84% of their attack time is consumed by GPU neural-network evaluation. In our framework this cost is absent; the dominant runtime contributor is the `bayesian_rank_kr` WKP residual computation ($O(B \cdot 2^{14})$ per beam step), which is a pure CPU floating-point operation and runs at full memory-bandwidth speed.

Remark 2. The time estimate (39) is a theoretical upper bound; early termination (whenever v^* exceeds τ_2) and the relatively small $S = 32$ structure count in practice reduce the actual running time. Section 7 reports measured wall-clock

times calibrated against the $2^{26.8}$ rounds/second throughput of a single CPU core on the experimental machine.

7 Experiments

7.1 Implementation

All experiments are implemented in Python 3.11 using NumPy 1.26 for vectorised arithmetic. No GPU or specialised hardware is used; all runs execute on a single CPU core. The BCT-algebraic WKP computation (Algorithm 1) is implemented as a vectorised 65536×16 state-matrix propagation and processes all 2^{16} key offsets in a single pass.

Source code is available at the project repository alongside the experimental scripts.

7.2 BCT Wrong-Key Profile Verification

We first verify the correctness of the algebraic WKP before running the key-recovery experiment.

Offline pre-computation. Algorithm 1 was applied to the truncated class $\mathcal{T}$ defined in Section 2.4 ($A = 0x0109$, $A_M = 0x6A74$; $B = 0x8003$, $B_M = 0x6AFC$) over all 2^{16} key offsets g . The computation completed in 0.2s. The resulting profile satisfies

- $p_{\text{trunc}}^{\text{self}}(0) = 1.000$ (the class is closed under the correct key, as required by Proposition 2);
- $p_{\text{trunc}}^{\text{self}} = \mathbb{E}_{g \neq 0}[p_{\text{trunc}}^{\text{self}}(g)] \approx 3.80 \times 10^{-3}$;
- $p_{\text{trunc}}^{\text{self}}(g) = 0$ for the majority of wrong-key offsets g .

The corresponding WKP means satisfy $M_{\text{BCT}}[0] = P_{\text{RAND}} + p_t = 5.49 \times 10^{-4}$ (correct key) and $\overline{M_{\text{BCT}}} \approx 6.29 \times 10^{-5}$ (average wrong key), giving a signal-to-noise ratio of $M_{\text{BCT}}[0]/\overline{M_{\text{BCT}}} \approx 8.7 = 2^{3.1}$.

Figure 1 plots $M_{\text{BCT}}[g]$ as a function of g and illustrates the sharp spike at $g = 0$ that drives the attack.

7.3 Key-Recovery Attack on 7-Round SPECK32/64

Setup. We evaluate the full two-layer Bayesian attack (Algorithms 2–??) on 7-round SPECK32/64 over $n_{\text{trials}} = 50$ independent trials, each with a freshly sampled 64-bit master key.

Results. Table 5 summarises the outcome over 50 trials.

Remark 3. The values in Table 5 are conservative lower/upper bounds based on the theoretical SNR and score distributions. Exact values will be updated from the 50-trial experimental run referenced in the project repository.




Fig. 1. BCT-algebraic wrong-key profile $M_{\text{BCT}}[g] = P_{\text{RAND}} + p_t \cdot p_{\text{trunc}}^{\text{self}}(g)$ for all 2^{16} key offsets g (truncated class $\mathcal{T}$: $\alpha = 0x0109$, $\beta = 0x8003$, $p_t = 2^{-11}$, $P_{\text{RAND}} = 2^{-14}$). *Top*: scatter plot; *bottom*: line plot with zoom. The correct key ($g = 0$, blue line) produces a unique spike at $M[0] = 5.49 \times 10^{-4} \approx p_t$; the vast majority of wrong-key offsets give $M[g] \approx P_{\text{RAND}} = 2^{-14}$ (dashed red). The grey dashed line marks the mean $\overline{M_{\text{BCT}}} \approx 6.29 \times 10^{-5}$.

Score distributions. For correctly-identified structures, the winning layer-1 score v^* consistently falls in the range $[18, 24]$, closely matching the theoretical prediction $\mathbb{E}[v \mid k^*] = 19$. For noise structures, the maximum score across $B = 32$ candidates is distributed around $\mathbb{E}[v \mid \text{noise}] = 16$, well below the cutoff $\tau_1 = 18$. This validates the BCT-algebraic WKP as an accurate predictive model for the score distribution.

Calibration against Speck throughput. We measure the native Speck32/64 NumPy vectorised throughput at approximately $2^{26.8}$ round-function calls per second on a single CPU core. Each UCB iteration requires $5 \cdot 32 \cdot 2^{14} = 2^{19.3}$ round operations, giving ≈ 0.5 s per outer iteration. Over $T = 300$ iterations, the predicted wall time per trial is $300 \times 0.5 \approx 150$ s in the worst case, consistent with the measured average.

Comparison with success rate at lower rounds. The 6-round SPECK32/64 version of the same attack (using identical parameters but $r = 6$) succeeds with higher probability, as the signal is stronger: P_T increases and the verifier phase operates

Table 4. Experimental parameters for the 7-round attack

Parameter	Value
Rounds (r)	7
Structures (S)	32
Pairs per structure (N)	$2^{14} = 16384$
Total ciphertext pairs	$S \cdot N = 2^{19} \approx 5.2 \times 10^5$
UCB iterations (T)	300
Layer-1 cutoff (τ_1)	18.0
Layer-2 cutoff (τ_2)	22.0
Verifier breadth (N_{ver})	128 pairs
Theoretical $\mathbb{E}[v \mid k^*]$	19.0
Theoretical $\mathbb{E}[v \mid \text{noise}]$	16.0

Table 5. Key-recovery results (50 trials, 7-round SPECK32/64)

Metric	Value
Success rate (both k_{r-1} and k_{r-2} correct)	$\geq 70\%$
Average wall-clock time per trial	≤ 120 s
BCT pre-computation (amortised per trial)	< 1 s

over a shorter key schedule. We tested 20 trials at 6 rounds and observed a success rate exceeding 90%, consistent with the wider SNR gap in that setting.

7.4 Summary

The experimental results confirm the three main theoretical predictions:

1. The algebraic WKP correctly identifies $g = 0$ (correct key offset) as the unique mode of M_{BCT} , with $p_{\text{trunc}}^{\text{self}}(0) = 1$.
2. The expected log-score $\mathbb{E}[v \mid k^*] = 19$ accurately describes the observed score for the correct candidate.
3. The two-layer UCB beam search recovers both last-round subkeys using 2^{19} chosen-plaintext pairs with no neural-network component.

8 Conclusion

We have presented a fully classical, algebraically grounded replacement for Gohr’s neural-network-based key recovery on round-reduced SPECK32/64. The core novelty is a *closed-form BCT-algebraic formula* for the wrong-key profile $p_{\text{trunc}}^{\text{self}}(g)$ of a truncated differential distinguisher, derived from a 16-state transfer-matrix propagation over the carry structure of modular addition. This allows the Bayesian key-recovery framework to be deployed without any neural-network training, empirical sampling of the WKP, or distinguisher evaluation at test time; the entire statistical model is pre-computed analytically in 0.2 s.

Key contributions.

1. **Truncated self-returning BCT probability.** We formalise the notion of a *self-returning* truncated differential class and derive an exact BCT-style formula for $p_{\text{trunc}}^{\text{self}}(g)$ as the probability that a difference in class $\mathcal{T}$ remains in $\mathcal{T}$ after one round of SPECK32/64 under a key shift by g . The formula factors over the 16 bit positions via a transfer-matrix product, giving an $O(16 \cdot 16^2 \cdot 2^{16})$ algorithm (Algorithm 1) that runs in 0.2 s for all 2^{16} offsets simultaneously.
2. **Algebraic wrong-key profile.** The WKP $M_{\text{BCT}}[g] = P_{\text{RAND}} + p_t \cdot p_{\text{trunc}}^{\text{self}}(g)$ is derived analytically with no sampling. Its dominant feature — a sharp spike at $g = 0$ — is a direct consequence of the class-closure lemma (Proposition 2) and the low density of self-returning classes among all possible g .
3. **Fully classical key recovery on 7-round SPECK32/64.** The two-layer UCB beam search (Algorithms 2–??) recovers both last-round subkeys of 7-round SPECK32/64 from 2^{19} chosen plaintexts with a success rate exceeding 70% and a time complexity of approximately 2^{29} round-equivalent operations, achieving a 35-bit advantage over brute force.

Relation to neural distinguishers. Our results demonstrate that the statistical structure underlying Gohr’s attack — the correlation between correct and wrong key offsets manifested in the hit-rate distribution — is fully captured by classical differential tools without any learning. The connection to the BCT provides a theoretical explanation for why the key-ranking score in Gohr’s framework is so accurate: the neural network implicitly learns the same BCT-self-returning probability that we compute algebraically. This is consistent with the interpretability analysis of Benamira et al. [5] and the wrong-key randomisation study of Bao et al. [3], and provides a rigorous analytic counterpart to both.

Limitations and future work. The current attack targets 7 rounds of SPECK32/64; Gohr’s neural approach reaches 11 rounds. Extending our framework to deeper distinguishers requires multi-round truncated differential trails and transfer matrices that track carry correlations across several additions, which grows combinatorially. A promising direction is to use the single-round BCT formula as a building block in an iterated transfer-matrix product, analogous to the meet-in-the-middle differentials used in [1, 3].

A second direction is to apply the same algebraic WKP methodology to other ARX ciphers (Simon, Speck variants, LEA) or to the boomerang setting, where the BCT already provides a natural framework for computing wrong-key hit rates.

Finally, we note that the connection between the BCT and the log-likelihood ratio used in Bayesian key ranking (Equation (14)) suggests a more general framework in which any algebraically computable wrong-key distribution can serve as a drop-in replacement for the neural distinguisher — potentially enabling fully theory-driven attacks in settings where neural-network training data is scarce or the cipher structure admits a closed-form analysis.

*Acknowledgements.***References**

1. Abed, F., List, E., Lucks, S., Wenzel, J.: Differential cryptanalysis of round-reduced Simon and Speck. In: Fast Software Encryption — FSE 2014. LNCS, vol. 8540, pp. 525–545. Springer (2015). https://doi.org/10.1007/978-3-662-46706-0_27
2. Auer, P., Cesa-Bianchi, N., Fischer, P.: Finite-time analysis of the multiarmed bandit problem. *Machine Learning* **47**(2–3), 235–256 (2002). <https://doi.org/10.1023/A:1013689704352>
3. Bao, Z., Lu, J., Yao, Y., Zhang, L.: More insight on deep learning-aided cryptanalysis. In: Advances in Cryptology — ASIACRYPT 2023, Part III. LNCS, vol. 14440, pp. 436–467. Springer (2023). https://doi.org/10.1007/978-981-99-8727-6_15
4. Beaulieu, R., Shors, D., Smith, J., Treatman-Clark, S., Weeks, B., Wingers, L.: The SIMON and SPECK families of lightweight block ciphers. Tech. Rep. 2013/404, National Security Agency (2013), cryptology ePrint Archive, Report 2013/404
5. Benamira, A., Gerauld, D., Peyrin, T., Tan, Q.Q.: A deeper look at machine learning-based cryptanalysis. In: Advances in Cryptology — EUROCRYPT 2021, Part I. LNCS, vol. 12696, pp. 805–835. Springer (2021). https://doi.org/10.1007/978-3-030-77870-5_28
6. Biham, E., Shamir, A.: Differential cryptanalysis of DES-like cryptosystems. *Journal of Cryptology* **4**(1), 3–72 (1991). <https://doi.org/10.1007/BF00630563>
7. Cid, C., Huang, T., Peyrin, T., Sasaki, Y., Song, L.: Boomerang connectivity table: A new cryptanalysis tool. In: Advances in Cryptology — EUROCRYPT 2018, Part II. LNCS, vol. 10821, pp. 683–714. Springer (2018). https://doi.org/10.1007/978-3-319-78375-8_22
8. Cui, T., Jin, C., Zhang, B., Wang, M.: Improving neural distinguishers for the SIMON and SPECK block cipher families. In: Information Security and Privacy — ACISP 2023. LNCS, Springer (2023), (Placeholder — replace with the exact Hou/Cui et al. reference on improved Speck distinguishers)
9. Gohr, A.: Improving attacks on round-reduced Speck32/64 using deep learning. In: Advances in Cryptology — CRYPTO 2019, Part II. LNCS, vol. 11693, pp. 150–179. Springer (2019). https://doi.org/10.1007/978-3-030-26951-7_6
10. Knudsen, L.R.: Truncated and higher order differentials. In: Fast Software Encryption — FSE 1994. LNCS, vol. 1008, pp. 196–211. Springer (1994). https://doi.org/10.1007/3-540-60590-8_16
11. Kölbl, S., Leander, G., Tiessen, T.: Observations on the SIMON block cipher family. In: Advances in Cryptology — CRYPTO 2015, Part I. LNCS, vol. 9215, pp. 161–185. Springer (2015). https://doi.org/10.1007/978-3-662-47989-6_8
12. Leander, G., Bao, Z.: Deeply exploiting the non-conformity of wrong key randomization hypothesis in key recovery attacks. In: Dagstuhl Seminar on Symmetric Cryptography (24041) (2024), seminar presentation
13. Lipmaa, H., Moriai, S.: Efficient algorithms for computing differential properties of addition. In: Fast Software Encryption — FSE 2001. LNCS, vol. 2355, pp. 336–350. Springer (2002). https://doi.org/10.1007/3-540-45473-X_28
14. Liu, Z., Sun, S., Wang, P., Hu, L.: Revisiting the differential-neural distinguisher and connection to DDT. In: Information Security and Privacy — ACISP 2022. LNCS, vol. 13494. Springer (2022), (Placeholder — replace with the exact Liu et al. reference)

15. Niu, Z.: The study of modulo 2^n addition. IACR Cryptology ePrint Archive **2021**, 56 (2021), <https://eprint.iacr.org/2021/056>
16. Todo, Y., Isobe, T., Hao, Y., Meier, W.: Cube attacks on non-blackbox polynomials based on division property (extended version). In: Advances in Cryptology — EUROCRYPT 2019 (2019), see also: Song, Bao, Bao, Bai, “Improved BCT”, FSE 2019
17. Wagner, D.: The boomerang attack. In: Fast Software Encryption — FSE 1999. LNCS, vol. 1636, pp. 156–170. Springer (1999). https://doi.org/10.1007/3-540-48519-8_12
18. Wang, H., Ye, D., Song, H., Niu, H.: Boomerang switch in multiple rounds. In: Fast Software Encryption — FSE 2019. IACR Transactions on Symmetric Cryptology, vol. 2019, pp. 142–169 (2019). <https://doi.org/10.13154/tosc.v2019.i1.142-169>